

Stories of Your Life as Others: A Round-Trip Evaluation of LLM-Generated Life Stories Conditioned on Rich Psychometric Profiles

Anonymous Authors

Abstract

Language models are increasingly used as simulated participants, personas, and synthetic respondents in social-science research. Yet most personality evaluations test shallow surfaces: whether a model can answer questionnaires, follow short trait labels, or reproduce demographic response patterns. Open-ended narrative is a stricter test because it requires a profile to be expressed through episodes, values, explanations, and emotional stance rather than through explicit trait words. We ask whether measured psychometric profiles can survive this transformation. Using 290 participants from a social-interaction dataset with self-report psychometrics, we convert each participant’s measured profile into controlled prose, generate Life Story Interview-style first-person answers, and ask independent models to recover the original scores from the generated text alone. Generated 24-turn life-story responses preserve strong individual-difference signal: independent scorers recover HEXACO traits at $r = 0.76\text{--}0.79$ and broader trust, psychopathy, social-anxiety, and social-value-orientation targets at lower but meaningful levels. Additional tests rule out several simple shortcuts, find signal in base/pretrained generators, and provide bounded human-reader and participant-language references. Generated narrative can preserve measurable psychometric structure, but this evaluation measures profile-to-language preservation rather than whether synthetic life stories replace human life-story interviews.

1 Introduction

Personality is one of the central cases where language and individual differences meet. Traits are robust predictors of life outcomes (Roberts et al., 2007; Ozer & Benet-Martínez, 2006), and trait-relevant information appears in word choice, emotional stance, interpersonal style, self-explanation, and the kinds of episodes people choose to narrate. This connection motivates a long tradition of personality prediction from text, from the lexical hypothesis (Saucier & Goldberg, 1996; Boyd & Pennebaker, 2017) to modern computational models trained on interviews, essays, social media, and self-narratives (Fast & Funder, 2008; Dai et al., 2022; Oltmanns et al., 2026; Speer et al., 2026; Wright et al., 2026).

Large language models make the problem newly bidirectional. They are not only used to infer traits from language; they are asked to generate language from a target profile. Recent work on virtual personas suggests that extended interview-style backstories can bind models more strongly to target response distributions than thin demographic or biographical prompts (Kang et al., 2025), while work on cross-scale psychometric prediction suggests that models often compress quantitative trait inputs into natural-language summaries before reasoning from them (Liu et al., 2026). Together, these lines point to a central possibility: language may be more than an output surface. It may be a medium through which models organise, carry, and recover individual-difference structure.

Existing evaluations often leave that possibility under-tested. A model can endorse questionnaire items, repeat trait adjectives, or follow a style cue without preserving a coherent profile through open-ended text. We therefore introduce a round-trip evaluation of profile-to-narrative preservation. The input is a measured profile from a real participant in PARSEL

(Hrkalovic et al., 2025); the output is generated Life Story Interview-style text, using the LSI format as a structured way to elicit long-form first-person narrative (McAdams & McLean, 2013). Independent scorers then receive only the generated text and infer the source psychometric profile. The experiment asks whether measured individual differences remain attached to generated language across two transformations: scores to controlled profile prose, and profile prose to generated life-story narrative.

The study uses two main profile formats: *Psychometric Narrative Profile*, which renders measured psychometrics as prose without personal facts or explicit score language, and *Psychometric Narrative + Facts*, which adds non-diagnostic biographical anchors to the same profile. We generate 4-part and 24-turn LSIs across four main generators and reverse-score every main condition with five independent scorers. Additional appendix sweeps cover broader generator lanes, less-filtered variants, and a diffusion-language model. The contribution is a behavioural test of psychometric encoding and decoding through generated narrative text, not a claim that synthetic LSIs are ecologically equivalent to human interviews.

2 Related Work

Personality prediction from text. Inferring personality from language has a long history in computational linguistics and psychology. The lexical hypothesis posits that personality-relevant distinctions become encoded in natural language (Saucier & Goldberg, 1996), and computational work has used social media, essays, interviews, and spoken or written narratives to recover personality from human-authored text (Boyd & Pennebaker, 2017; Fast & Funder, 2008; Dai et al., 2022; Oltmanns et al., 2026; Speer et al., 2026; Wright et al., 2026). Reported rates vary with genre and dose: Dai et al. (2022) fine-tune InterviewBERT on interview responses, Oltmanns et al. (2026) use McAdams-style narratives, Wright et al. (2026) use zero-shot LLM scoring of stream-of-thought writing, and Speer et al. (2026) compare NLP methods on open-ended responses. These studies establish the recognition side of the problem: personality can be inferred from open human text, though rates depend on genre, length, prompt structure, and scoring method. Our setting reverses the direction. We start from measured psychometrics, generate LSI-style text, and ask what survives the transformation.

LLM persona conditioning. Recent work asks whether language models can adopt stable personas or simulate human response distributions when conditioned on profile information (Argyle et al., 2023; Park et al., 2023; Bai et al., 2025; Serapio-García et al., 2025; Shi et al., 2025; Sorokovikova et al., 2024; Zhou et al., 2025). Most such work uses invented personas, demographic skeletons, thin trait descriptions, or questionnaire self-report by the conditioned model; these choices matter because questionnaire self-report can diverge from task behaviour (Han et al., 2025), and socially grounded persona frameworks require richer information than a trait label alone (Venkit et al., 2026). Kang et al. demonstrate that extended McAdams-style backstories produce deeper persona binding than short descriptions (Kang et al., 2025), while work on cross-scale psychometric prediction finds that models often translate quantitative traits into natural-language summaries before reasoning from them (Liu et al., 2026). These results motivate a stricter generation-and-recovery test of whether measured psychometric profiles bind more strongly when rendered into controlled narrative language.

Life-story interviews and narrative personality. The McAdams Life Story Interview elicits personality through narrative themes rather than direct trait self-description (McAdams, 2001; McAdams & McLean, 2013). Its structure of chapters, key scenes, challenges, values, and future scripts is useful here as a structured prompt format for long-form first-person narrative. Human LSI and self-narration studies indicate that narrative text can carry personality signal (Lodi-Smith et al., 2009; Oltmanns et al., 2026; Wright et al., 2026). At the same time, LLM processing of life narratives can introduce representational shifts and bias (Subbiah et al., 2026). We therefore use the LSI format to structure generated first-person narrative and evaluate recoverable psychometric structure, while reserving matched human-LSI validity for future data collection.

Post-training and mechanism. Work on assistant personas and post-training suggests that chat alignment can privilege a default assistant mode and change behavioural surrogacy (Lu et al., 2026; Binz et al., 2026). Work on model personality and representation also illustrates why mechanism remains difficult: behavioural capability, corpus exposure, and internal representation are related but separable questions (Ju et al., 2025). We therefore pair instruction-tuned systems with base/pretrained controls and OLMo-family corpus inspection, while treating full acquisition provenance as future mechanistic work.

These literatures leave four empirical questions for the present paper:

- RQ1.** Does measured psychometric structure survive encoding into generated life-story narrative?
- RQ2.** Is recovery robust to controls against direct score handoff, factual grounding, provider dialect, transparent trait vocabulary, and scorer artefacts?
- RQ3.** Does recovery appear in base/pretrained generators, or is it created by instruction tuning?
- RQ4.** Do human readers and short participant-authored conversations provide bounded evidence that generated stories carry person-profile signal?

3 Method

3.1 Data and psychometric targets

We used PARSEL (Hrkalovic et al., 2025), a multimodal dataset collected to study social perception and partner selection in cooperative settings. PARSEL includes 297 participants with psychometric questionnaires, pairwise conversations, and collaborative tasks. We used the $N = 290$ participants with complete profiles for the round-trip analyses. A subset of 248 participants had sufficient conversation data for the human-language reference analyses.

Each participant completed the HEXACO-60 inventory (Ashton & Lee, 2009), yielding Honesty-Humility, Emotionality, Extraversion, Agreeableness, Conscientiousness, and Openness. We also used ten non-HEXACO continuous targets: four trust subscales (Propensity, Ability, Integrity, Benevolence; Mayer & Davis 1999), four Psychopathic Personality Traits Scale subscales (Cognitive Responsiveness, Affective Responsiveness, Egocentricity, Interpersonal Manipulation; Boduszek et al. 2016), the Social Interaction Anxiety Scale (Mattick & Clarke, 1998), and social value orientation. We summarised recovery over three aggregates: HEXACO6, Beyond10, and Continuous16.

3.2 Round-trip pipeline

Profile construction. The first transformation encoded measured scores into natural-language profile text. *Psychometric Narrative Profile* was the clean psychometric condition: it rendered measured scores as prose while excluding personal facts and forbidding explicit instrument names, exact scores, percentiles, construct identifiers, and scoring language where possible. *Psychometric Narrative + Facts* added non-diagnostic biographical anchors to the same psychometric narrative profile. *Facts-Only Control* contained biographical facts without psychometric interpretation and was used as a control rather than as a main treatment condition. Appendix C reports the full profile-writer bake-off and rationale for selecting Gemma 4 31B IT as the default writer.

The profile-writing stage was audited directly rather than treated as hidden prompting infrastructure. Gemma 4 31B IT was selected as the default profile writer after a full- N profile-writer comparison because it preserved strong direct decodability while remaining locally hostable. Appendix C reports the full direct-scoring table for this first transformation.

LSI generation. The second transformation expanded the profile into first-person responses using an adapted LSI prompt format (McAdams & McLean, 2013). The main panel crossed four generators (Gemini 3 Flash, GPT-4.1, Gemma 4 31B IT, and Qwen 3.6 27B), two profile conditions (*PsychNarr* and *PsychNarr+Facts*), and two generation lengths (4-part and 24-turn

LSIs). Four-part LSIs tested whether compact narratives already carry recoverable signal; 24-turn LSIs tested the richer long-form setting most aligned with the LSI interview format. All main rows used $N = 290$. Appendix D reports additional 4-part generator sweeps, including smaller, open-weight, less-filtered, and diffusion-language model lanes.

Blind psychometric scoring. Following recent zero-shot LLM scoring work on personality from open text (Wright et al., 2026), the decoding arm received only the generated LSI text and inferred the source psychometrics. The main scoring panel used Gemini 3 Flash, Sonnet 4.6, GPT-5.4-Mini, Qwen 3.6 Plus, and Gemma 4 31B IT. Scores were compared with the source profiles using Pearson correlations and then averaged over HEXACO6, Beyond10, and Continuous16 targets. Appendix E reports scoring prompts, parsing checks, human-reader calibration, and the PARSEL conversation coding structure.

3.3 Validation and reference analyses

We used several control families to ask what carried the recoverable signal. Direct scorecard controls tested whether exact numeric handoff or construct glossaries recovered the same signal as narrative profile text. Facts-only controls tested whether biographical anchors alone carried psychometric geometry. Family-bias checks compared same-family generator/scorer pairs against cross-family pairs. Additional generator sweeps asked whether the result was confined to a narrow autoregressive chat-model ecology. Transparent schema tests asked how much recovery could be explained by obvious trait-language lexicons. Severe sentence-drop ablations deleted every generated sentence containing any transparent-schema lexical hit before rescoring. Human-reader calibration asked whether non-LLM readers could match short person-description snippets to generated LSI theme snippets above chance.

We also included two reference analyses. True base/pretrained generators tested whether recoverable psychometric-language structure appeared before instruction tuning. PARSEL conversations provided the only participant-authored language tied to the same source profiles. These conversations are short task interactions, so they offered human-language context rather than format-matched validation.

4 Results

We report results in the order of the four research questions. First, we test whether measured psychometric structure survives the full round trip from source profile to generated LSI text and back to inferred scores. Second, we test alternative explanations: direct score handoff, biographical facts, provider-specific generator/scorer effects, and transparent trait vocabulary. Third, we compare true base/pretrained generators with matched post-trained or instruct counterparts to ask whether recovery depends on instruction tuning. Fourth, we report two bounded human-language references: a non-LLM reader matching task on generated snippets and short PARSEL task conversations from the same source participants.

4.1 RQ1: Round-trip recovery across generated LSIs

The first question was whether measured psychometric structure survives the two-step transformation from scores to profile prose to generated life-story narrative. We first scored the profile text directly to estimate how much source structure was available before LSI generation. We then scored the generated LSIs to test how much of that structure survived narrative expansion. The profile-writing transformation was highly informative on its own: Gemini 3 Flash direct scoring of the profile text reached HEXACO/Beyond10/Continuous16 values of 0.92/0.78/0.83 for *PsychNarr* and 0.91/0.78/0.83 for *PsychNarr+Facts*. The generated-LSI matrix tested the second transformation: whether independent scorers could recover the source psychometrics from generated life-story text alone. It crossed 4 generators, 2 profile conditions, 2 LSI lengths, and 5 scorers, for 80 full- N rows.

Generated LSIs preserved strong recoverable signal (Table 1; Figure 1). For 24-turn LSIs, *PsychNarr* reached mean HEXACO6/Beyond10/Continuous16 recovery of 0.79/0.62/0.68; *PsychNarr+Facts* reached 0.76/0.61/0.67. The 4-part LSIs were also strong, with Continu-

ous16 recovery of 0.66 for *PsychNarr* and 0.65 for *PsychNarr+Facts*. Longer LSIs improved recovery modestly, but the effect was already visible in compact generated text.

The recovery pattern was not confined to one scorer or generator. Mean Continuous16 recovery across all main rows was 0.68 for Sonnet 4.6, 0.67 for Gemini 3 Flash, 0.67 for Qwen 3.6 Plus, 0.66 for Gemma 4 31B IT, and 0.64 for GPT-5.4-Mini. Gemini 3 Flash and Gemma 4 31B IT were strongest as generators, GPT-4.1 remained close, and Qwen 3.6 27B was weaker but still well above null.

Table 1: Round-trip recovery averaged across four LSI generators and five independent scorers.

Condition	LSI length	Rows	HEXACO6	Beyond10	Continuous16
<i>PsychNarr</i>	4-part	20	0.76	0.60	0.66
<i>PsychNarr+Facts</i>	4-part	20	0.73	0.60	0.65
<i>PsychNarr</i>	24-turn	20	0.79	0.62	0.68
<i>PsychNarr+Facts</i>	24-turn	20	0.76	0.61	0.67




Figure 1: Generated 24-turn LSIs preserve recoverable psychometric signal across both profile formats. Bars report mean Pearson r across generators and scorers; error bars report 95% confidence intervals over generator-scorer rows.

4.2 RQ2: Controls against handoff, facts, dialect, and transparent schema

The second question was whether recovery could be explained by simpler routes. Direct numeric handoff carried signal, as it should: a compact exact-score scorecard recovered HEXACO/Beyond10/Continuous16 at 0.65/0.52/0.57, and a longer glossary scorecard recovered 0.73/0.58/0.64 (Table 2). Both scorecard controls recovered less signal than the matched *PsychNarr* reference, especially for Continuous16 and SVO.

Facts-only controls gave the complementary result. Generated facts-only inputs recovered little HEXACO signal: 0.10 for Gemini-generated facts-only controls and 0.15 for Gemma-generated facts-only controls. Facts-only rows therefore ruled out biographical anchors as the primary source of recovery in *PsychNarr+Facts*. Same-family generator/scorer cells also did not outperform cross-family cells (HEXACO/Beyond10 0.75/0.60 versus 0.76/0.61), so the effect was not higher when the same model family generated and scored the text.

Table 2: Controls separating narrative profile conditioning from direct score handoff and facts-only inputs.

Control	HEXACO6	Beyond10	Continuous16	SVO
Matched <i>PsychNarr</i> reference	0.76	0.62	0.67	0.57
Compact exact-score scorecard	0.65	0.52	0.57	0.41
Glossary scorecard	0.73	0.58	0.64	0.55
Gemini facts-only control	0.10	–	–	–
Gemma facts-only control	0.15	–	–	–

Transparent trait-language controls recovered part of the HEXACO signal. A lexical schema built from public trait and facet descriptors, not model outputs, recovered mean HEXACO6 signal of 0.51 without an LLM judge. After residualising LLM recovery on all six HEXACO lexicons, mean HEXACO6 recovery remained 0.67; after residualising on every schema lexicon plus length, it remained 0.64, with all 192/192 tests significant. A more severe deletion test removed every generated sentence containing any transparent-schema lexical hit before rescoring. For full- N Gemini 3 Flash 24-turn LSIs, this retained only 37.0% of words for *PsychNarr* and 39.9% for *PsychNarr+Facts*, yet HEXACO6 recovery remained 0.75 and 0.73 respectively. OLMo sentence-drop controls followed the same pattern. Appendix H reports the lexicon construction, residualisation, sentence-drop tables, and schema-control figure.

4.3 RQ3: Base/pretrained generators and post-training

RQ3 asked whether recovery depends on instruction tuning. To test this, we compared four true base/pretrained generators with matched post-trained or instruct counterparts where a close sibling model was available.

Four true base/pretrained generators recovered HEXACO signal at full $N = 290$: Mistral Small 24B, Gemma 4 31B, Qwen3.5-35B-A3B, and OLMo 3 32B (Figure 2; Appendix G). Matched post-trained or instruct 4-part LSI runs were higher in all four families under the available scorer coverage: Mistral rose from 0.45–0.46 to 0.71, Gemma from 0.47–0.52 to 0.74–0.78, Qwen from 0.65–0.67 to 0.81–0.83, and OLMo from 0.49–0.53 to 0.65.




Figure 2: Base/pretrained generators recover HEXACO signal, while matched post-trained or instruct counterparts recover more. Bars report family-level HEXACO recovery; error bars report the available independent-scorer ranges.

4.4 RQ4: Human-reader matching and short-conversation references

The fourth question was whether any evidence outside LLM reverse scoring supports the bounded interpretation that generated stories carry person-profile signal. We first ran a small non-LLM reader calibration. Twenty-five native-English Prolific raters each completed 25 four-option matches between short person-description snippets and generated LSI theme snippets. They selected the correct match on 372/625 trials, for mean accuracy of 59.5% against a 25% chance rate ($t(24) = 8.24$, one-tailed $p = 9.38 \times 10^{-9}$). The task was deliberately lightweight, but it indicates that non-LLM readers can recover person-profile signal from limited generated materials. Appendix E reports the human-reader design and the human-language coding structure used for PARSEL conversations.

The PARSEL conversation references were modest. Participants’ own HEXACO scores predicted coded features of their short task conversations at mean $|r| = 0.09$, with a featurewise best-trait reference of 0.18. Partner-perceived HEXACO after conversation

correlated with self-report at mean $r = 0.16$, and LLM scoring of the short conversations recovered HEXACO at $r = 0.18$ – 0.22 . Against these low-bandwidth references, generated-story-to-real-conversation same-feature correspondence was $r = 0.15$ – 0.18 . These values provide the denominator for interpreting the generated-LSI-to-conversation same-feature correlations.




Figure 3: Short PARSEL conversations provide a low-bandwidth participant-language reference. Generated-LSI-to-conversation correspondence falls in the same range as other short-conversation psychometric references, while generated-LSI round-trip recovery is much higher because it uses the full generated narrative.

As a secondary content-level check, profile-conditioned LSIs also had broader emotional-valence variability than unconditioned baselines, and the Sentimentality-to-valence-variability effect replicated across independent valence coders. Appendix F reports the full emotional-reactivity analysis.

5 Discussion

Measured psychometric structure survives transformation into generated life-story narrative. Across the main 24-turn conditions, independent scorers recovered HEXACO traits at $r = 0.76$ – 0.79 and recovered broader individual-difference targets at lower but consistent levels. The simplest shortcuts did not match the full narrative condition: exact scorecards, facts-only controls, same-family generator/scorer pairs, and transparent trait lexicons all recovered less signal than the primary LSI rows. Base/pretrained generators also recovered signal, indicating that instruction tuning is not required, although matched post-trained models performed better. Finally, human readers and short PARSEL conversations provide bounded external references: they support the claim that generated materials carry person-profile signal, while leaving matched human-LSI validation for future work.

The central finding is profile-to-narrative preservation in a strictly operational sense: source individual-difference structure remains attached to generated language across transformations. The round trip is therefore a measurement design for two linked capabilities. The first capability is encoding measured psychometrics into profile and life-story language. The second is decoding psychometrics back from that language. The effect is robust across multiple generators, multiple scorers, two profile formats, two LSI lengths, direct scorecard controls, facts-only controls, transparent-schema tests, severe sentence deletion, human-reader calibration, true base/pretrained generators, and exploratory generator sweeps that include a diffusion-language model.

The narrative transformation matters because it gives the model an intermediate representation richer than a score table. Compact exact-score scorecards and longer glossary scorecards recover meaningful signal, but they do not match the narrative profile condition.

Facts-only controls recover little HEXACO signal. The controlled profile prose can express covariance, tension, and interaction among traits: affective warmth can be separated from anxiety, trust in benevolence can be distinguished from trust in ability, and interpersonal caution can coexist with conscientious self-control. The LSI generator then expands that person-language into episodes, explanations, values, emotional responses, and interpersonal scenes. This interpretation fits adjacent evidence that LLMs can use natural-language summaries as potent psychometric compressions (Liu et al., 2026), while extending it into a stricter generation-and-recovery setting.

The base-model comparison constrains where the capability can be placed. The base/pretrained controls indicate that recoverable psychometric-language structure is already available before assistant-style instruction tuning, while matched post-trained/instruct comparisons indicate that post-training can amplify task performance. This pattern complements Kang et al.’s persona-binding result (Kang et al., 2025): deep persona methods are not merely about asking a chat model to role-play, but about activating representational resources that can already be present in pretrained models. It also matters because post-training can reposition models toward a default assistant mode (Lu et al., 2026) and, in some behavioural-surrogate settings, reduce human-response alignment (Binz et al., 2026). Our result makes the profile-to-language mapping a useful probe for future work on how pretraining and post-training shape social-cognitive capabilities.

The schema results should be read constructively. A fine-grained psychometric schema is a plausible mechanism: models may learn rich associations among traits, facets, affective style, interpersonal warmth, anxiety, trust, and language. That possibility does not invalidate the finding. A schema rich enough to predict Sentimentality rather than Anxiety, affective-responsiveness deficits rather than unrelated dark-trait facets, and Benevolence rather than unrelated trust facets is preserving real psychometric structure. The transparent-schema tests indicate that obvious lexical versions of this explanation are incomplete. Richer latent schemas, broader personality-language mappings, and ordinary social regularities in pretraining remain targets for future mechanistic and corpus work.

The participant-language references are deliberately modest. Human open-ended narratives and life narrative interviews can carry recoverable personality signal (Oltmanns et al., 2026; Wright et al., 2026), but those studies use richer and more introspective text than the available PARSEL conversations. Wright et al. report LLM-self-report convergence of $r = 0.44\text{--}0.48$ from 30-minute stream-of-thought writing and multi-day diaries; the strict PARSEL conversation reference aggregates about 12.7 minutes and 1,110 target-speaker words per participant in short dyadic tasks. We therefore treat the conversation result as a low-bandwidth participant-language reference. It indicates that some generated LSI features fall in the same small range as other short-conversation psychometric references, but it does not validate synthetic LSIs as substitutes for human life-story interviews. The secondary emotional-reactivity analysis in Appendix F provides a complementary content-level check that profiles shape narrative dynamics, not only final score recovery.

Kang et al.’s work on deep binding combines behavioural persona tests with corpus checks, not a full provenance account (Kang et al., 2025). Our base-model comparisons and OLMo-family corpus probes belong in that empirical tradition: they narrow and stress-test the interpretation, while leaving mechanistic acquisition as future work. The stronger validation study would collect matched human LSIs, long-form self-narration, or reflective essays from the same participants.

6 Limitations

This study tests psychometric-linguistic encoding and decoding through generated profile prompts and LSI-style text. It does not establish that synthetic LSIs faithfully reproduce how the same people would write or speak in matched human LSIs, and it does not establish unrestricted person simulation. The evidence is English-only and drawn from PARSEL’s Western, educated, industrialised, rich, and democratic (WEIRD) participant context, so we do not know whether the same profile-to-language preservation would survive translation, multilingual profile construction and scoring, or lower-resource linguistic contexts.

The pipeline is also LLM-dense: language models write profiles, generate LSIs, score psychometrics, and support several analyses. Cross-family scoring, deterministic controls, transparent lexicons, sentence deletion, base models, and human-reader calibration reduce shared-method concerns, but they do not eliminate them.

We also do not identify the training-data source or mechanistic origin of the mapping. OLMo-family corpus tools can inspect ingredients in pretraining data, and base-model controls move part of the phenomenon before instruction tuning, but how language models generalise from training data into social-cognitive capabilities is a much larger open problem. Mapping latent space may reveal the capability’s structure; reconstructing exactly how it came to be is beyond this paper. The human-reader pilot is a calibration on generated snippets, not expert annotation of original PARSEL transcripts, and the conversation reference uses short task conversations with modest recoverable psychometric signal. The next validation step is a matched dataset with psychometrics, human LSIs or long-form self-narration, participant-authored conversations, and consent structures designed for both human and synthetic narrative analysis.

7 Conclusion

Generated life-story narratives can preserve measured psychometric structure across a profile-to-language-to-score round trip. Starting from PARSEL participants’ self-report psychometrics, we converted each measured profile into controlled prose, generated LSI-style first-person narratives, and withheld the source profile from independent scorers. Those scorers recovered strong HEXACO signal in the main 24-turn conditions ($r = 0.76\text{--}0.79$), with lower but meaningful recovery for broader trust, psychopathy, social-anxiety, and social-value-orientation targets. Recovery remained visible across generators, scorers, profile formats, LSI lengths, direct scorecard controls, facts-only controls, transparent-schema tests, sentence deletion, human-reader calibration, and true base/pretrained generator controls. Together, these results indicate that generated narrative can carry measurable individual-difference structure. They also define the next validation problem: testing how that generated structure compares with human life-story interviews from the same people.

This finding raises the standard for persona and simulated-participant evaluation. The current state of the art often tests whether models endorse questionnaires, follow trait labels, reproduce demographic patterns, or remain consistent with short persona descriptions. Those tests are useful, but they leave open whether a profile survives when it must be expressed through open-ended language. Narrative makes that question harder. It forces profile information to pass through episodes, explanations, values, emotional responses, and interpersonal scenes rather than remaining a labelled score vector. For social-science simulation and behavioural model evaluation, this means that profile adoption should be tested in the language where simulated people are expected to act.

The next step is to connect this round-trip test to human narrative evidence and model mechanism. Matched human LSIs, long-form self-narration, multilingual samples, and richer participant-authored conversations would clarify how generated narratives compare with human accounts from the same people. Mechanistic and corpus-level analyses would clarify how models represent, transform, and sometimes distort profile information before and after instruction tuning. The contribution of this study is to make profile-to-narrative preservation measurable, so future work can ask not only whether a model can simulate a profile, but where that profile lives in generated language and how faithfully it survives transformation.

Reproducibility Statement

All psychometric data are from PARSEL (Hrkalovic et al., 2025). Model versions, prompts, scoring parameters, statistical analysis scripts, and derived results are staged in the accompanying repository. The manuscript reports values from the consolidated camera-ready artefact package, with table and figure provenance documented in the appendix map and evidence manifests.

Ethics Statement

This study uses existing psychometric data collected under ethical approval. No directly identifiable information appears in the manuscript. *Psychometric Narrative + Facts* profiles include source-person life anchors for experimental conditioning; we therefore treat them as controlled research artefacts and distinguish them from *Psychometric Narrative Profile*, which removes personal facts. The capacity to recover personality information from generated text raises privacy considerations that warrant careful attention in deployment contexts.

AI Disclosure Statement

As the generative pipeline is itself the subject of this research, large language models served as experimental instruments throughout: as profile writers, narrative generators, personality scorers, content coders, matchers, leakage auditors, and corpus-inspection aids. The human authors were additionally aided in literature search, coding support, statistical checking, and manuscript preparation by generative agents operating within customised coding harnesses. Generated narratives were evaluated for psychometric signal recovery and behavioural correspondence, not subjective quality alone. All methodological decisions, experimental design choices, and interpretive claims are by the human authors, who are responsible for the accuracy, academic integrity, and intellectual contributions of this paper.

References

- Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023. doi: 10.1017/pan.2023.2.
- Michael C. Ashton and Kibeom Lee. The HEXACO-60: A short measure of the major dimensions of personality. *Journal of Personality Assessment*, 91(4):340–345, 2009.
- Yuqi Bai, Tianyu Huang, Kun Sun, and Yuting Chen. Scaling law in llm simulated personality: More detailed and realistic persona profile is all you need, 2025. URL <https://arxiv.org/abs/2510.11734>.
- Marcel Binz, Elif Akata, Abdullah Almaatouq, Mohammed Alsobay, Oleksii Ariasov, Franziska Brändle, David Broska, Jason W. Burton, Nuno Busch, Frederick Callaway, Vanessa Cheung, Brian Christian, et al. Post-training makes large language models less human-like, 2026. URL <https://arxiv.org/abs/2605.07632>.
- Daniel Boduszek, Agata Debowska, Katie Dhingra, and Matt DeLisi. Introduction and validation of psychopathic personality traits scale (ppts) in a large prison sample. *Journal of Criminal Justice*, 46:9–17, 2016. ISSN 0047-2352. doi: <https://doi.org/10.1016/j.jcrimjus.2016.02.004>. URL <https://www.sciencedirect.com/science/article/pii/S0047235216300046>.
- Ryan L Boyd and James W Pennebaker. Language-based personality: a new approach to personality in a digital world. *Current Opinion in Behavioral Sciences*, 18:63–68, 2017. ISSN 2352-1546. doi: <https://doi.org/10.1016/j.cobeha.2017.07.017>. URL <https://www.sciencedirect.com/science/article/pii/S2352154617300487>. Big data in the behavioural sciences.
- Yimeng Dai, Madhura Jayaratne, and Buddhi Jayatilleke. Explainable personality prediction using answers to open-ended interview questions. *Frontiers in Psychology*, 13:865841, 2022. doi: 10.3389/fpsyg.2022.865841.
- Lisa A. Fast and David C. Funder. Personality as manifest in word use: correlations with self-report, acquaintance report, and behavior. *Journal of Personality and Social Psychology*, 94(2):334–346, February 2008. ISSN 0022-3514. doi: 10.1037/0022-3514.94.2.334.
- Pengrui Han, Rafal Kocielnik, Peiyang Song, Ramit Debnath, Dean Mobbs, Anima Anandkumar, and R. Michael Alvarez. The personality illusion: Revealing dissociation between self-reports & behavior in llms, 2025. URL <https://arxiv.org/abs/2509.03730>.
- Tiffany Matej Hrkalo, Bernd Dudzik, Daniel Balliet, and Hayley Hung. PARSEL: A multimodal dataset for modeling decision-making processes involved in selecting partners for joint tasks. *IEEE Transactions on Affective Computing*, 16(4):3481–3498, 2025. doi: 10.1109/TAFFC.2025.3600687.
- Yiwen Ju et al. Probing then editing response personality of large language models. In *Proceedings of COLM*, 2025.
- Minwoo Kang, Suhong Moon, Seung Hyeon Lee, Ayush Raj, Joseph Suh, and David Chan. Deep binding of language model virtual personas: a study on approximating political partisan misperceptions. In *Second Conference on Language Modeling*, 2025. URL <https://openreview.net/forum?id=zHdSctNmM4>.
- Yi-Fei Liu, Yi-Long Lu, Di He, and Hang Zhang. From five dimensions to many: Large language models as precise and interpretable psychological profilers, 2026. URL <https://arxiv.org/abs/2511.03235>.
- Jennifer Lodi-Smith, Aaron C. Geise, Brent W. Roberts, and Richard W. Robins. Narrating personality change. *Journal of Personality and Social Psychology*, 96(3):679–689, March 2009. ISSN 0022-3514. doi: 10.1037/a0014611.

- Christina Lu, Jack Gallagher, Jonathan Michala, Kyle Fish, and Jack Lindsey. The assistant axis: Situating and stabilizing the default persona of language models, 2026. URL <https://arxiv.org/abs/2601.10387>.
- Richard P. Mattick and J.Christopher Clarke. Development and validation of measures of social phobia scrutiny fear and social interaction anxiety. *Behaviour Research and Therapy*, 36(4):455–470, 1998. ISSN 0005-7967. doi: [https://doi.org/10.1016/S0005-7967\(97\)10031-6](https://doi.org/10.1016/S0005-7967(97)10031-6). URL <https://www.sciencedirect.com/science/article/pii/S0005796797100316>.
- Roger C. Mayer and James H. Davis. The effect of the performance appraisal system on trust for management: A field quasi-experiment. *Journal of Applied Psychology*, 84(1): 123–136, February 1999. ISSN 1939-1854, 0021-9010. doi: [10.1037/0021-9010.84.1.123](https://doi.org/10.1037/0021-9010.84.1.123). URL <https://doi.apa.org/doi/10.1037/0021-9010.84.1.123>.
- Dan P. McAdams. The psychology of life stories. In *Review of General Psychology*, volume 5, pp. 100–122. SAGE, 2001.
- Dan P. McAdams and Kate C. McLean. Narrative identity. *Current Directions in Psychological Science*, 22(3):233–238, June 2013. ISSN 0963-7214, 1467-8721. doi: [10.1177/0963721413475622](https://doi.org/10.1177/0963721413475622). URL <https://journals.sagepub.com/doi/10.1177/0963721413475622>.
- Joshua R. Oltmanns, Ritik Khandelwal, Jerry Ma, Jocelyn Brickman, Tu Do, Rasiq Hussain, and Mehak Gupta. Language-based ai modeling of personality traits and pathology from life narrative interviews. *Journal of Psychopathology and Clinical Science*, 135(1):122–135, January 2026. ISSN 2769-755X. doi: [10.1037/abn0001047](https://doi.org/10.1037/abn0001047).
- Daniel J. Ozer and Verónica Benet-Martínez. Personality and the prediction of consequential outcomes. *Annual Review of Psychology*, 57:401–421, 2006. ISSN 0066-4308. doi: [10.1146/annurev.psych.57.102904.190127](https://doi.org/10.1146/annurev.psych.57.102904.190127).
- Joon Sung Park, Joseph O’Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. 2023. doi: [10.1145/3586183.3606763](https://doi.org/10.1145/3586183.3606763). URL <https://doi.org/10.1145/3586183.3606763>.
- Brent W. Roberts, Nathan R. Kuncel, Rebecca Shiner, Avshalom Caspi, and Lewis R. Goldberg. The power of personality: The comparative validity of personality traits, socioeconomic status, and cognitive ability for predicting important life outcomes. *Perspectives on Psychological Science*, 2(4):313–345, 2007. doi: [10.1111/j.1745-6916.2007.00047.x](https://doi.org/10.1111/j.1745-6916.2007.00047.x). URL <https://doi.org/10.1111/j.1745-6916.2007.00047.x>. PMID: 26151971.
- Gerard Saucier and Lewis R. Goldberg. The language of personality: Lexical perspectives on the five-factor model. In *The five-factor model of personality: Theoretical perspectives*, pp. 21–50. Guilford Press, New York, NY, US, 1996. ISBN 9781572300682.
- Gregory Serapio-García, Mustafa Safdari, Clément Crepy, Luning Sun, Stephen Fitz, Peter Romero, Marwa Abdulhai, Aleksandra Faust, and Maja Matarić. A psychometric framework for evaluating and shaping personality traits in large language models. *Nature Machine Intelligence*, 7(12):1954–1968, 2025. ISSN 2522-5839. doi: [10.1038/s42256-025-01115-6](https://doi.org/10.1038/s42256-025-01115-6).
- Quan Shi, Carlos E Jimenez, Stephen Dong, Brian Seo, Caden Yao, Adam Kelch, and Karthik R Narasimhan. IMPersona: Evaluating individual level LLM impersonation. In *Second Conference on Language Modeling*, 2025. URL <https://openreview.net/forum?id=7qhBXqONLN>.
- Aleksandra Sorokovikova, Sharwin Rezaghali, Natalia Fedorova, and Ivan P. Yamshchikov. LLMs simulate big5 personality traits: Further evidence. In Ameet Deshpande, EunJeong Hwang, Vishvak Murahari, Joon Sung Park, Diyi Yang, Ashish Sabharwal, Karthik Narasimhan, and Ashwin Kalyan (eds.), *Proceedings of the 1st Workshop on Personalization of Generative AI Systems (PERSONALIZE 2024)*, pp. 83–87, St. Julians, Malta, March 2024. Association for Computational Linguistics. doi: [10.18653/v1/2024.personalize-1.7](https://doi.org/10.18653/v1/2024.personalize-1.7). URL <https://aclanthology.org/2024.personalize-1.7/>.

- Andrew B. Speer, Angie Y. Delacruz, Takudzwa A. Chawota, James Perrotta, and Cort W. Rudolph. Unpacking the validity of open-ended personality assessments using fine-tuned large language models. *Journal of Organizational Research Methods (ORM)*, pp. 10944281251413746, March 2026. ISSN 1094-4281, 1552-7425. doi: 10.1177/10944281251413746. URL <https://journals.sagepub.com/doi/10.1177/10944281251413746>.
- M. Subbiah, H. Mian, N. Deas, A. Mayukha, Dan P. McAdams, and Kathleen McKeown. Whose story gets told? positionality and bias in llm summaries of life narratives, 2026. URL <https://arxiv.org/abs/2604.20131>.
- Pranav Narayanan Venkit, Yu Li, Yada Pruksachatkun, and Chien-Sheng Wu. The need for a socially-grounded persona framework for user simulation. 2026. URL <https://arxiv.org/abs/2601.07110>.
- Aidan G. C. Wright, Whitney R. Ringwald, Colin E. Vize, Johannes C. Eichstaedt, Mike Angstadt, Aman Taxali, and Chandra Sripada. Assessing personality using zero-shot generative ai scoring of brief open-ended text. *Nature Human Behaviour*, January 2026. ISSN 2397-3374. doi: 10.1038/s41562-025-02389-x.
- Lingfeng Zhou, Jialing Zhang, Jin Gao, Mohan Jiang, and Dequan Wang. Personaeval: Are LLM evaluators human enough to judge role-play? In *Second Conference on Language Modeling*, 2025. URL <https://openreview.net/forum?id=drdrFhKYjP>.

A Appendix Roadmap

This appendix is an evidence map for the camera-ready version. The main text reports the compact argument; the appendix records the controls, implementation routes, and interpretation boundaries needed to audit the result. The accompanying repository contains the full prompts, generated outputs, scorer rows, parser outputs, and statistical scripts.

Appendix section	Purpose	Key evidence route
Profile conditions	Defines the public condition names and maps them to internal run labels.	Appendix B.
Profile writer bake-off	Documents that the score-to-profile transformation was selected empirically, not treated as hidden prompt craft.	Appendix C; full- <i>N</i> profile-writer rows scored by Gemini 3 Flash.
Model sweep	Documents the central 80-row matrix, broader 4-part generator lanes, less-filtered variants, and diffusion-language-model check.	Appendix D; consolidated run manifests.
Scoring and coding	Documents reverse scoring, parsing checks, human-reader calibration, and PARSEL conversation denominators.	Appendix E.
Emotional reactivity	Reports the secondary content-level analysis of valence variability and independent valence coders.	Appendix F.
Base and post-trained generators	Separates pretraining-available capability from post-training amplification.	Appendix G; full- <i>N</i> base/pretrained and matched instruct/post-trained rows.
Schema and sentence deletion	Tests whether obvious trait-language schemas explain recovery.	Appendix H; lexicon construction, residualisation, and sentence-drop outputs.
Corpus inspection	Adapts Kang-style corpus checks and OLMo-Trace/InfiniGram probes to generated LSI outputs and schema-hit sentences.	Appendix I; <code>corpus_inspection_bundle_20260608/README.md</code> .

Table 3: Camera-ready appendix map. The table is a routing surface for the larger evidence forest, not a replacement for the linked manifests and analysis outputs.

B Profile Conditions and Terminology

The paper uses public-facing condition names in the main text while preserving older internal run labels in the artefact tree. The public names are deliberately descriptive.

Public condition	Short label	Meaning
<i>Psychometric Narrative Profile</i>	<i>PsychNarr</i>	A controlled prose profile constructed from measured psychometrics. It excludes personal facts and avoids explicit instrument names, exact scores, percentiles, and scoring language where possible.
<i>Psychometric Narrative + Facts</i>	<i>PsychNarr+Facts</i>	The same psychometric narrative profile, augmented with non-diagnostic biographical anchors. Older artefacts sometimes call this “FullBio” or “ArtWeave”.
<i>Facts-Only Control</i>	<i>FactsOnly</i>	Biographical anchors without psychometric interpretation. This is a control, not a main treatment condition.
Scorecard controls	Scorecard	Deterministic, non-narrative source-data handoffs: compact exact-score rows and a longer glossary variant.

Table 4: Condition names used in the camera-ready paper.

C Profile Writer Bake-Off

The main experiments use Gemma 4 31B IT to write the psychometric narrative profiles. This choice followed a full-*N* comparison of candidate profile writers. The strongest direct profile decode came from GPT-5.5, but its *PsychNarr+Facts* profiles were much longer. Gemma 4 31B IT preserved high direct decodability while producing compact profiles and remaining open-weight and locally hostable.

Profile writer/condition	N	Mean words	Median words	HEXACO	Beyond10	PPTS total
Gemma 4 31B IT <i>PsychNarr</i>	290	930	929	0.92	0.78	0.87
Gemma 4 31B IT <i>PsychNarr+Facts</i>	290	1115	1114	0.91	0.78	0.87
Opus 4.7 <i>PsychNarr</i>	290	1060	1059	0.92	0.74	0.77
Opus 4.7 <i>PsychNarr+Facts</i>	290	1395	1389	0.90	0.72	0.77
Qwen 3.6 <i>PsychNarr</i>	290	1192	1186	0.91	0.74	0.82
Qwen 3.6 <i>PsychNarr+Facts</i>	290	1627	1520	0.89	0.74	0.82
Qwen 3.7 <i>PsychNarr</i>	290	1111	1106	0.92	0.74	0.78
Qwen 3.7 <i>PsychNarr+Facts</i>	290	1603	1580	0.91	0.73	0.79
GPT-5.5 <i>PsychNarr</i>	290	1524	1522	0.92	0.80	0.88
GPT-5.5 <i>PsychNarr+Facts</i>	290	2940	2932	0.92	0.79	0.88

Table 5: Full- N profile-writer comparison scored by Gemini 3 Flash. These rows evaluate the first transformation directly: measured psychometrics to controlled profile language.

D Model Sweep and Generation Variants

The main matrix is the 80-row panel summarised in Table 1: four LSI generators, two profile formats, two LSI lengths, and five independent scorers. All rows use $N = 290$. The main generators are Gemini 3 Flash, GPT-4.1, Gemma 4 31B IT, and Qwen 3.6 27B; the main scorers are Gemini 3 Flash, Sonnet 4.6, GPT-5.4-Mini, Qwen 3.6 Plus, and Gemma 4 31B IT.

Additional 4-part sweeps were used to stress-test the ecology around that main panel. These include smaller and open-weight generators, safety-filtered and less-filtered (“Heretic”) variants, and Mercury, a diffusion language model. The generator sweep is not used to inflate the headline; it is used to check whether recovery is an artefact of one narrow provider or one autoregressive chat-model style.

Sweep family	Question	Summary
Main 80-row matrix	Does recovery survive model/scorer variation?	Four generators, two profile formats, two LSI lengths, five scorers, all full- N . Same-family generator/scorer cells do not outperform cross-family cells on HEXACO/Beyond10 (0.75/0.60 vs. 0.76/0.61).
Broader 4-part lanes	Is the result limited to the four main generators?	Additional local and API lanes include smaller, open-weight, and alternative provider models. These are appendix robustness checks rather than headline rows.
Less-filtered variants	Is ordinary safety filtering suppressing or manufacturing recovery?	Full- N <i>PsychNarr+Facts</i> 4-part comparison changed HEXACO6 by approximately zero for Gemma and +0.04 for Qwen across four judges.
Diffusion-language model	Is recovery specific to autoregressive generation?	The broader sweep includes Mercury as a diffusion-language-model lane; this is treated as exploratory evidence against a narrow autoregressive artefact.

Table 6: Generation-variant checks supporting the main matrix.

E Scoring, Human Reader Calibration, and Conversation Coding

The main reverse-scoring panel uses five independent LLM scorers. Each scorer receives only the generated LSI text and predicts the source psychometric targets. Scoring follows the zero-shot open-text personality-scoring logic of Wright et al. (2026), adapted to generated LSI-style responses and to the HEXACO plus beyond-HEXACO target set. Parsed scorer outputs were checked for row count, participant identifiers, missing targets, and out-of-range values before correlations were computed.

The human-reader calibration asks a narrower question: can non-LLM readers detect person-profile signal in short generated materials? Twenty-five native-English Prolific raters completed 625 four-option matches between short person-description snippets and generated LSI theme snippets. They selected the correct match on 372 trials, for mean accuracy of 59.5% against a 25% chance rate ($t(24) = 8.24$, one-tailed $p = 9.38 \times 10^{-9}$). This is not matched human-LSI validation; it is a calibration that the generated materials carry person-profile signal visible to human readers.

Table 7: Human-reader calibration on short snippets of generated materials.

Item	Value
Native-English Prolific raters	25
Four-option matches completed	625
Correct matches	372
Mean accuracy	59.5%
Chance accuracy	25%
One-sample test	$t(24) = 8.24, p = 9.38 \times 10^{-9}$
Attention checks	75/75 passed
Mean completion time	60.6 minutes

The PARSEL conversation analyses use the only participant-authored language tied to the same source profiles. Conversations are short cooperative-task interactions, so the reference is intentionally low-bandwidth. Participants’ own HEXACO scores predict coded conversation features at mean $|r| = 0.09$, with a featurewise best-trait reference of 0.18. Partner-perceived HEXACO after conversation correlates with self-report at mean $r = 0.16$, and LLM scoring of the short conversations recovers HEXACO at $r = 0.18$ –0.22. These denominators keep the generated-LSI-to-conversation bridge in proportion.

Table 8: PARSEL conversation references. Conversations are short task interactions, so they provide low-bandwidth human-language context.

Reference	Pearson r
Source HEXACO $\rightarrow$ coded conversation features	0.09
Featurewise best-trait reference	0.18
Partner-perceived HEXACO $\rightarrow$ self-report	0.16
LLM scoring of short conversations	0.18–0.22
Generated LSI $\rightarrow$ conversation same-feature bridge	0.15–0.18

F Generated Emotional Reactivity

The emotional-reactivity analysis tests whether profile conditioning affects generated narrative dynamics, not only final psychometric scores. Profile-conditioned LSIs vary more in emotional-valence variability than unconditioned LSIs: conditioned lanes have between-story SDs of 0.10–0.13, while unconditioned baselines range from 0.07–0.09 (Figure 4). Brown-Forsythe tests are significant in 27/32 conditioned-versus-unconditioned comparisons, while unconditioned baselines do not differ from each other.

The Sentimentality-to-valence-variability effect also replicates across independent valence coders, with mean partial correlations of 0.45 for Qwen 3.7 Max, 0.45 for Grok 4.20, and 0.42 for Xiaomi MiMo. These content-level checks are secondary to the main round-trip recovery result, but they indicate that the profile shapes generated narrative dynamics as well as scorer-level trait estimates.




Figure 4: Unconditioned generated narratives have narrower emotional-valence variability than profile-conditioned LSI lanes.

G Base/Pretrained and Post-Trained Generators

Reviewer discussion led us to expand the base/pretrained analysis to a full four-family panel, then to add matched post-trained or instruct 4-part counterparts. The central interpretation is qualitative as well as numerical: the capability is not created by instruction-tuned chat compliance, but post-training makes the task stronger and easier to control.

Table 9: Matched base/pretrained and post-trained/instruct 4-part LSI controls. Values are HEXACO6 recovery; ranges reflect available independent scorers and available scorer coverage differs slightly by family.

Model family	Base/pretrained	Post-trained/instruct
Mistral Small 24B	0.45–0.46	0.71
Gemma 4 31B	0.47–0.52	0.74–0.78
Qwen3.5-35B-A3B	0.65–0.67	0.81–0.83
OLMo 3 32B	0.49–0.53	0.65

The base/pretrained rows use true base or pretrained checkpoints for Mistral Small 24B, Gemma 4 31B, Qwen3.5-35B-A3B, and OLMo 3 32B. The matched post-trained/instruct values use the closest available sibling rows under available scorer coverage. Scorer parity is strongest for the Qwen family; Mistral and OLMo instruct values are Gemini-scored, and Gemma IT values come from the main 4-part *PsychNarr+Facts* scorer matrix.

H Transparent Schema and Sentence-Drop Tests

Transparent lexical schemas were built from public trait and facet descriptors rather than from model outputs. They intentionally operationalise the simplest version of a psychometric-schema account: if obvious public trait language carries the result, a non-LLM lexical scoring scheme should recover much of the same signal. It does recover a real part of the signal. Across eight headline 24-turn lanes, the transparent schema alone recovers mean HEXACO6 signal of 0.51. But the LLM-scoring result remains substantial after residualising on the schema scores: raw HEXACO6 recovery averages 0.77, partial recovery after all six HEXACO lexicons is 0.67, and partial recovery after every schema lexicon plus length is 0.64, with 192/192 tests significant.

The sentence-drop control is deliberately more severe than literal overlap deletion. Every generated sentence containing any transparent-schema lexical hit is removed before rescoreing. For Gemini 3 Flash 24-turn LSIs, this leaves no transparent schema hits and retains only 37.0% of words for *PsychNarr* and 39.9% for *PsychNarr+Facts*. Recovery survives. For OLMo base LSIs, deletion retains 56.1% of words; for OLMo instruct LSIs, deletion retains 43.4% of words. These controls do not rule out a rich latent schema. They indicate that the obvious lexical version is incomplete.

Table 10: Transparent-schema controls. Sentence-drop rows delete every sentence containing any transparent-schema lexical hit before rescoreing.

Test	HEXACO6	Note
Raw LLM scoring, 8 main 24-turn lanes	0.77	Original recovery
Transparent lexical schema alone	0.51	No LLM judge
Partial after all six HEXACO lexicons	0.67	Residualised
Partial after all schema lexicons + length	0.64	Residualised, 192/192 significant
Gemini <i>PsychNarr</i> sentence drop	0.81 → 0.75	37.0% words retained
Gemini <i>PsychNarr+Facts</i> sentence drop	0.79 → 0.73	39.9% words retained
OLMo base sentence drop	0.53 → 0.45	56.1% words retained
OLMo instruct sentence drop	0.65 → 0.57	43.4% words retained

We also ran non-LLM lexical-valence checks using Empath-style lexical categories on base-model outputs. These analyses are secondary because Empath is blunt relative to the LLM scorers, but they help separate the content result from LLM-only scoring. The staged outputs live under `corpus_inspection_bundle_20260608/14_base_empath_variability/`.




Figure 5: Transparent-schema tests. Left: obvious lexical schemas recover part of the signal, but substantial recovery remains after residualising on schema lexicons. Right: recovery remains after deleting every sentence containing any transparent-schema lexical hit.

I Corpus Inspection Bundle

The OLMo inspection work lives in `corpus_inspection_bundle_20260608/`. The bundle uses symlinks to preserve provenance to the original analysis outputs while providing a compact retrieval surface. Its front door is `corpus_inspection_bundle_20260608/README.md`; its itemised map is `corpus_inspection_bundle_20260608/MANIFEST.tsv`. The purpose is bounded corpus inspection, not provenance proof.

I.1 Kang-style generated-output n-grams

Following the spirit of Kang et al.’s generated-output corpus appendix, we computed generated 5-grams and 10-grams from four generated-LSI corpora and queried exact counts against two OLMo-family InfiniGram indexes. No raw PARSEL conversations or source profile prompts were submitted to InfiniGram; queries use generated LSI surface strings only.

Generated corpus	Files	Tokens	10-grams	Notes
Gemini 24-turn <i>PsychNarr</i>	290	2,934,047	2,931,437	Main psychometric narrative profile lane.
Gemini 24-turn <i>Psych-Narr+Facts</i>	290	2,836,886	2,834,276	Main psychometric narrative plus facts lane.
OLMo base 4-part	290	836,708	834,098	Full-N OLMo base lane used for closed-loop probes.
OLMo instruct 4-part	290	414,906	412,296	Full-N OLMo instruct lane assembled from 40-PID and remaining-250 runs.

Table 11: Generated corpora used for Kang-style n-gram inspection.

Short 5-grams often appear in OLMo-family indexes because they are ordinary narrative phrases. Generated 10-grams are much less often exact pretraining strings. Among the top

20 generated 10-gram rows per corpus, zero exact hits were found in both checked OLMo-family indexes for 18/20 *PsychNarr* rows, 19/20 *PsychNarr+Facts* rows, 12/20 OLMo base rows, and 20/20 OLMo instruct rows. Anchor-conditioned 10-grams had the same pattern: 32/33, 32/33, 26/33, and 31/33 rows respectively had zero exact hits in both checked indexes. The full tables are in `analysis/kang_style_corpus_appendix_20260608/`.

I.2 Psychometric bridge phrase counts

Exact-count bridge probes found individual ingredients in the checked OLMo-family indexes, but not obvious recipe phrases linking the instruments and LSI format.

Exact phrase	Count
HEXACO	22,615
Life Story Interview	453
personal narrative	509,630
Psychopathic Personality Traits Scale	135
Social Interaction Anxiety Scale	1,335
Social Value Orientation	1,584
HEXACO Life Story Interview	0
HEXACO personal narrative	0
HEXACO Psychopathic Personality Traits Scale	0
HEXACO Trust benevolence	0

Table 12: Representative exact-count probes over checked OLMo-family indexes. Counts support an “ingredients without obvious recipe” interpretation, not an absence-of-evidence claim over all possible training text.

Covariance-literature probes had a similar split. In the checked exact phrase set, explicit bridge phrases produced 0/30 nonzero exact phrases (total count 0), ingredient phrases produced 22/32 nonzero exact phrases (total count 32,410), and literature-handle phrases produced 0/10 nonzero exact phrases. The full query set is staged under `corpus_inspection_bundle_20260608/04_covariance_literature_probe/`.

I.3 OLMoTrace closed-loop checks

We then traced the sentences most favourable to the transparent-schema concern: deleted sentences from OLMo-family outputs that contained transparent schema hits. For OLMo base, 40/40 endpoint calls succeeded and returned 1,550 source rows. Returned sources were diffuse: uncategorised web text (31.9%), education/jobs (20.9%), health (16.1%), literature (10.1%), finance/business (4.8%), art/design (4.2%), politics (3.5%), and science/technology (3.4%). Returned matched spans did not contain explicit psychometric bridge language. For OLMo instruct, 40/40 endpoint calls succeeded and returned 1,547 source rows; a strict scan found zero explicit psychometric labels or bridge handles in returned span text or source URIs.

These traces are deliberately modest. They indicate that the trait-bearing deleted sentences most favourable to a transparent-schema account trace to broad ordinary web, life, affect, health, education, literary, and topical sources rather than obvious psychometric bridge recipes. They do not identify how OLMo learned the capability, rule out all relevant corpus examples, or locate internal features.

I.4 Corpus-inspection boundary

The corpus-inspection bundle can support four bounded conclusions. First, generated LSI text uses ordinary narrative ingredients. Second, common generated longer strings are generally not verbatim pretraining copies in the checked OLMo-family indexes. Third, obvious exact psychometric-to-LSI bridge phrases and covariance recipe phrases were not found in the checked exact-count probes, even though individual ingredients were present.

Fourth, the OLMo-family sentence-drop controls indicate that recovery survives deletion of the most transparent trait-language sentences at full N .

It cannot establish the exact training-data source of the psychometric-language mapping, prove absence of all relevant examples, or make mechanistic claims about circuits, features, attention heads, or latent activation geometry. Those are future work. The present paper uses corpus inspection the way the main method uses behavioural controls: to bound simpler explanations while keeping the acquisition question open.

J Evidence Manifest

The consolidated rebuttal/camera-ready evidence forest is larger than the manuscript can reproduce. The following local routes are the intended audit surfaces for camera-ready revision and post-acceptance cleanup.

Route	Contents
00_front_door/LOAD_BEARING_INDEX.md	Front-door index for the rebuttal evidence package, including main runs, controls, and analysis outputs.
00_front_door/FINAL_TABLE_STAGING_20260525.md	Older table staging surface used to reconcile submitted-paper claims with CCV1 results.
00_front_door/NUMBER_TO_ARTIFACT_MANIFEST_20260525.md	Number-to-artifact route map for rebuttal figures, tables, and checks.
tables/ and figures/camera_ready/	Camera-ready table snippets and rebuilt main figures.
analysis/kang_style_corpus_appendix_20260608/	Generated n-gram and InfiniGram exact-count outputs for corpus inspection.
corpus_inspection_bundle_20260608/	Compact public-facing bundle for OLMoTrace/InfiniGram, schema, sentence-drop, base-control, and covariance probes.

Table 13: Primary local audit routes for the camera-ready evidence package.